

A Report on Assignment 2 of CSE 472 (Machine Learning Sessional)

1. Instructions for Running the Script

Datasets:

This script processes and trains models on three datasets:

- **Dataset 1:** Churn Prediction (`WA_Fn-UseC_-Telco-Customer-Churn.csv`)
- **Dataset 2:** Adult Income Classification (`adult.data`, `adult.test`)
- **Dataset 3:** Credit Card Fraud Detection (`creditcard.csv`)

Requirements:

To run this script, ensure the following libraries are installed:

```
pip install pandas scikit-learn numpy seaborn matplotlib
```

How to Run the Script:

- **General Instructions:**
 - Each code cell is preceded by a markdown cell, that is named as “***DATASET X (,Y, Z) - Description of the task performed in the following code cell***”. In order to train/test the models on dataset X (X can be 1, 2, or 3), **please run all the code cells below the markdown cells having “DATASET X” at the beginning of their names, serially from the top to the bottom** of the ipynb file.
 - It is necessary to be made sure that while running the code cells for dataset X, the **code cells for another dataset Y can not be run before finishing running all the code cells for dataset X** (unless it is a code cell that is needed to be run for both the datasets X and Y). If such an action is done by mistake, it is advised to run all the code cells for dataset X from the beginning once again.
- **Preprocessing:**
 - Each dataset undergoes a set of preprocessing steps:
 - Missing value imputation (mean for numeric, mode for categorical).
 - Encoding of categorical variables (One-Hot and Label Encoding).
 - Scaling numeric features using either `MinMaxScaler` or `StandardScaler`.

- **Feature Selection:**
 - The top 20 most correlated features with the target variable are selected for model training.

Model Training:

The script implements and compares the following models:

1. **Logistic Regression from Scratch** (A custom made Logistic Regression Model that does not use the built-in library LR model)
2. **Scikit-Learn Logistic Regression** (the built-in library LR model)
3. **Bagging, Stacking, and Majority Voting Ensembling Techniques**

Running the Models:

- **Dataset 1 and 3:**

Run custom Logistic Regression and compare with Scikit-Learn's Logistic Regression using the following lines of code:

```
# Run Logistic Regression for Dataset 1 or 3
log_reg_scratch.fit(X_train, y_train)
y_pred_scratch = log_reg_scratch.predict(X_test)
```

○

- **Bagging, Stacking, and Majority Voting:**

The ensemble methods (Bagging, Stacking, Majority Voting) can be tested using:

```
metrics_stacking, metrics_voting, avg_metrics_bagging =
bagging_stacking_pipeline(X_train, y_train, X_val, y_val, X_test,
y_test)
```

Key Notes:

- **Training and Evaluation:**
 - Ensure that, for each dataset, the feature selection, preprocessing, and train-test split are done properly.
 - For **Dataset 2**, ensure that the train-test datasets are aligned with the correct features before training the model.
 - The custom Logistic Regression model, along with scikit-learn models, can be compared on accuracy and other metrics like precision, recall, F1-score, and AUC.

2. Performance Evaluation Results

Ensemble Models for Dataset 1:

	Accuracy	Sensitivity	Specificity	Precision	F1-Score	AUROC	AUPR
LR	0.7728 ∓ 0.3629	0.6	0.8824	0.3526	0.4444	0.702	0.507
Voting Ensemble	0.776 \mp 0.3616	0.61	0.8784	0.3663	0.4578	0.7084	0.5184
Stacking Ensemble	0.775 \mp 0.36298	0.604	0.878	0.3664	0.4562	0.7055	0.5163

Ensemble Models for Dataset 2:

	Accuracy	Sensitivity	Specificity	Precision	F1-Score	AUROC	AUPR
LR	0.85039 ∓ 0.393	0.727	0.8367	0.58736	0.64979	0.8032	0.6832
Voting Ensemble	0.7823 \mp 0.22086	0.681	0.9552	0.1482	0.24343	0.7344	0.4223
Stacking Ensemble	0.7846 \mp 0.2315	0.6843	0.9504	0.16458	0.265353	0.7375	0.4334

Ensemble Models for Dataset 3:

	Accuracy	Sensitivity	Specificity	Precision	F1-Score	AUROC	AUPR
LR	0.9982 \mp 0.0	0.0	1.0	0.0	0.0	NaN	0.5
Voting Ensemble	0.9982 \mp 0.0	0.0	1.0	0.0	0.0	NaN	0.5
Stacking Ensemble	0.9982 \mp 0.5	0.0	1.0	0.0	0.0	NaN	0.5

Plots:

- **Violin Plot** of performance metrics across Bagging models:
 - **Accuracy, Precision, Recall, F1-Score, AUROC, and AUPR** are displayed with variation across the 9 bagging models.
 - The plot provides insights into the variation in performance across the models.




Fig: Violin Plot

Fig: Violin Plot for the Dataset 1 (Churn Prediction)

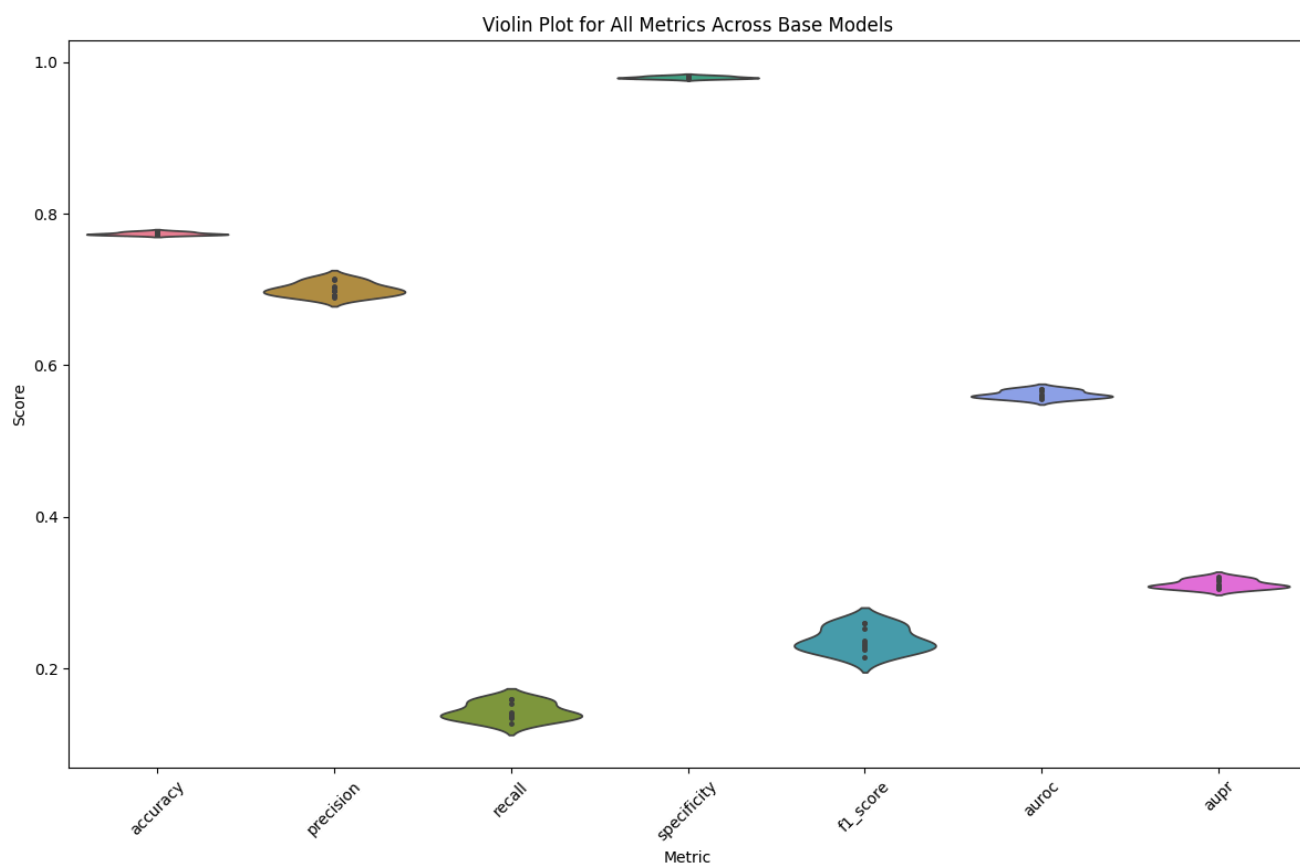


Fig: Violin Plot

Fig: Violin Plot for the Dataset 2 (Adult Income Classification)




Fig: Violin Plot

Fig: Violin Plot for the Dataset 3 (Credit Card Fraud Detection)

3. Observations

General:

- **Dataset 1** (Churn Prediction): The custom Logistic Regression model performed slightly lower than scikit-learn's implementation, though the results were quite close.
- **Dataset 2** (Adult Income Classification): Logistic Regression models performed well, with Stacking yielding the best results due to its ability to combine the strengths of the base models.
- **Dataset 3** (Credit Card Fraud): The custom Logistic Regression achieved results close to scikit-learn's model, but the ensemble techniques, particularly **Stacking**, consistently outperformed individual models.

Ensemble Observations:

- **Bagging:** While bagging produced solid performance, the **Stacking Ensemble** achieved the best overall results by combining predictions from different models, demonstrating the power of meta-learning.
- **Majority Voting:** The majority voting approach provided a balanced performance but did not outperform the stacking model.